

Gold Price Prediction

Importing the Libraries

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split
from sklearn.ensemble import RandomForestRegressor
from sklearn import metrics
```

Data Collection and Processing

```
# loading the csv data to a Pandas DataFrame
gold_data = pd.read_csv('/content/gld_price_data.csv')
```

```
# print first 5 rows in the dataframe
gold_data.head()
```

	Date	SPX	GLD	USO	SLV	EUR/USD
0	1/2/2008	1447.160034	84.860001	78.470001	15.180	1.471692
1	1/3/2008	1447.160034	85.570000	78.370003	15.285	1.474491
2	1/4/2008	1411.630005	85.129997	77.309998	15.167	1.475492
3	1/7/2008	1416.180054	84.769997	75.500000	15.053	1.468299
4	1/8/2008	1390.189941	86.779999	76.059998	15.590	1.557099

```
# print last 5 rows of the dataframe
gold_data.tail()
```

	Date	SPX	GLD	USO	SLV	EUR/USD
2285	5/8/2018	2671.919922	124.589996	14.0600	15.5100	1.186789
2286	5/9/2018	2697.790039	124.330002	14.3700	15.5300	1.184722
2287	5/10/2018	2723.070068	125.180000	14.4100	15.7400	1.191753
2288	5/14/2018	2730.129883	124.489998	14.3800	15.5600	1.193118
2289	5/16/2018	2725.780029	122.543800	14.4058	15.4542	1.182033

```
# number of rows and columns
gold_data.shape
```

```
(2290, 6)
```

```
# getting some basic informations about the data
gold_data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 2290 entries, 0 to 2289
Data columns (total 6 columns):
 #   Column      Non-Null Count  Dtype
---  -
 0   Date        2290 non-null   object
 1   SPX         2290 non-null   float64
 2   GLD         2290 non-null   float64
 3   USO         2290 non-null   float64
 4   SLV         2290 non-null   float64
 5   EUR/USD     2290 non-null   float64
dtypes: float64(5), object(1)
memory usage: 107.5+ KB
```

```
# checking the number of missing values
gold_data.isnull().sum()
```

```
Date      0
SPX        0
GLD        0
USO        0
SLV        0
EUR/USD    0
dtype: int64
```

```
# getting the statistical measures of the data
gold_data.describe()
```

	SPX	GLD	USO	SLV	EUR/USD
count	2290.000000	2290.000000	2290.000000	2290.000000	2290.000000
mean	1654.315776	122.732875	31.842221	20.084997	1.283653
std	519.111540	23.283346	19.523517	7.092566	0.131547
min	676.530029	70.000000	7.960000	8.850000	1.039047
25%	1239.874969	109.725000	14.380000	15.570000	1.171313
50%	1551.434998	120.580002	33.869999	17.268500	1.303296
75%	2073.010070	132.840004	37.827501	22.882499	1.369971
max	2872.870117	184.589996	117.480003	47.259998	1.598798


Correlation:

1. Positive Correlation
2. Negative Correlation

```
correlation = gold_data.corr()
```

```
# constructing a heatmap to understand the correlation
plt.figure(figsize = (8,8))
sns.heatmap(correlation, cbar=True, square=True, fmt='.1f',annot=True, annot_kws={'size':8})
```

<matplotlib.axes._subplots.AxesSubplot at 0x7fc6c1339390>




```
# correlation values of GLD
print(correlation['GLD'])
```

```
SPX      0.049345
GLD      1.000000
USO     -0.186360
SLV      0.866632
EUR/USD  -0.024375
Name: GLD, dtype: float64
```

```
# checking the distribution of the GLD Price
sns.distplot(gold_data['GLD'],color='green')
```

```
/usr/local/lib/python3.7/dist-packages/seaborn/distributions.py:2619: FutureWarning:
  warnings.warn(msg, FutureWarning)
<matplotlib.axes._subplots.AxesSubplot at 0x7fc6b8d1db50>
```



Splitting the Features and Target

```
X = gold_data.drop(['Date', 'GLD'],axis=1)
Y = gold_data['GLD']
```

```
print(X)
```

	SPX	USO	SLV	EUR/USD
0	1447.160034	78.470001	15.1800	1.471692
1	1447.160034	78.370003	15.2850	1.474491
2	1411.630005	77.309998	15.1670	1.475492
3	1416.180054	75.500000	15.0530	1.468299
4	1390.189941	76.059998	15.5900	1.557099
...
2285	2671.919922	14.060000	15.5100	1.186789
2286	2697.790039	14.370000	15.5300	1.184722
2287	2723.070068	14.410000	15.7400	1.191753
2288	2730.129883	14.380000	15.5600	1.193118
2289	2725.780029	14.405800	15.4542	1.182033

```
[2290 rows x 4 columns]
```

```
print(Y)
```

0	84.860001
1	85.570000
2	85.129997
3	84.769997
4	86.779999
...	...

```

2285    124.589996
2286    124.330002
2287    125.180000
2288    124.489998
2289    122.543800
Name: GLD, Length: 2290, dtype: float64

```

Splitting into Training data and Test Data

```
X_train, X_test, Y_train, Y_test = train_test_split(X, Y, test_size = 0.2, random_state=2)
```

Model Training: Random Forest Regressor

```
regressor = RandomForestRegressor(n_estimators=100)
```

```

# training the model
regressor.fit(X_train,Y_train)

RandomForestRegressor(bootstrap=True, ccp_alpha=0.0, criterion='mse',
                        max_depth=None, max_features='auto', max_leaf_nodes=None,
                        max_samples=None, min_impurity_decrease=0.0,
                        min_impurity_split=None, min_samples_leaf=1,
                        min_samples_split=2, min_weight_fraction_leaf=0.0,
                        n_estimators=100, n_jobs=None, oob_score=False,
                        random_state=None, verbose=0, warm_start=False)

```

Model Evaluation

```

# prediction on Test Data
test_data_prediction = regressor.predict(X_test)

```

```

print(test_data_prediction)

177.38329918 114.64039959 119.19960017  94.60850088 125.9105002
166.16170117 114.78830078 116.72750144  88.23249863 148.79410093
120.43709949  89.49189963 112.42740005 117.16410062 118.76730107
 88.26819957  94.17959991 117.03980027 118.68860187 120.46700061
126.69219821 121.93169977 151.0317001  164.8883009 118.57249957
120.45210147 149.27170066 118.41109917 172.1916982 106.11009901
104.91510102 149.02780119 113.93390085 124.77760132 147.57520008
119.66130115 115.33560056 112.7179001  113.59920193 143.72850106
118.06739759 102.8665003  115.95050147 103.96580166  98.65870041
117.61180062  90.71590012  91.59580046 153.54069904 102.7856998
155.15610086 114.2913015  139.59730133  90.08199819 115.45249941
114.16229967 123.18690026 121.65890068 165.23500136  92.82109957
134.85450148 121.37869913 120.94280086 104.61560026 141.97710287
121.82589949 116.7792005  113.36690074 126.96039798 122.68459933
125.75569952 121.24000017  86.87839896 132.24210114 146.27910094
 92.68879946 158.26489988 158.58370291 126.24269904 164.82119949
108.8909997 109.69600103 103.69829825  94.45410041 127.55860275
106.96670056 161.22579953 121.78870032 132.10340021 130.83690201
160.73099973  90.05979869 174.11350152 128.60310058 126.76329874
 86.36809921 124.42249893 150.20059705  89.70189993 107.08719946

```

```

108.97100004 85.13819853 135.97499947 154.48120263 139.21290345
74.71359996 152.49590117 125.86930029 126.69040057 127.46729924
108.5891995 156.02080028 114.43390151 117.01080169 125.39479912
154.13720149 121.34759985 156.33229894 92.89730063 125.49800129
125.6943005 87.82630023 92.0712993 126.23239924 128.8117044
113.05190065 117.64609735 121.04040024 127.02099811 120.19050056
136.63010135 93.88119911 119.84680048 113.28520079 94.16339945
108.97819954 87.4174988 108.96709951 89.4240999 92.42290038
131.82000302 162.36740104 89.36169991 119.56550052 133.29070178
124.04260032 128.55100206 101.89989836 89.00159859 132.01170093
120.07559995 108.28759999 168.94870105 115.2914005 86.63329889
118.64610046 91.14609965 161.52070042 116.74270042 121.50179962
160.35549832 120.15499921 112.75329906 108.44869881 126.71399985
76.2044001 103.01509992 127.38300249 121.8899988 92.58240017
132.79000111 118.0406011 116.22929972 154.10250311 158.65240054
109.92139951 157.20539775 119.29380075 159.90010168 118.4356004
157.60329976 115.03659929 116.73640027 149.9081986 114.90390075
125.73989882 165.48489944 117.47949998 124.99299923 153.28720384
153.37650249 132.10700054 114.84490043 121.15340214 125.02000073
89.79180031 123.22240011 154.72580153 111.56170019 106.68719988
162.68980167 118.34219971 165.62609875 134.12970126 115.07129979
153.05659894 168.78400083 115.12760042 114.03420112 158.71769947
85.21639922 127.0533007 127.91340015 128.82849954 124.37240057
124.10310042 90.69620029 153.3353999 96.90439992 137.83170015
89.0958994 107.70480008 115.1630006 112.6551009 124.44729896
91.43679872 125.50450135 162.28739942 120.0169986 165.07400098
126.73729813 112.25840014 127.6229994 94.96289867 91.15499989

103.44289909 120.89539999 83.0883994 126.34360017 160.04540488
117.36220109 118.4873999 119.83159971 122.79719982 120.11700122
121.58769967 118.19740048 107.13569998 148.29159994 126.46119828
115.81090131 74.10299996 127.8334011 154.2249012 122.7043002
125.64650059 88.77689998 103.1299985 124.82390015 120.27730032
73.33760061 151.69379971 121.19720067 104.6042 86.26739788
115.19279914 171.87689806 119.7112005 159.74459713 113.25219969
121.41780022 118.75260067 96.06619987 119.20199982 125.93160037
118.52489962 95.80870028 153.76010203 121.98980047 147.55079988
159.36440268 113.73570001 122.57969935 149.33399788 127.02320036
165.90130023 135.30230085 120.07469949 166.6752973 108.3125992
121.75459853 138.33960149 106.92659892]

```

```
# R squared error
```


```
error_score = metrics.r2_score(Y_test, test_data_prediction)
print("R squared error : ", error_score)
```

```
R squared error : 0.9893081607893282
```

Compare the Actual Values and Predicted Values in a Plot

```
Y_test = list(Y_test)
```

```
plt.plot(Y_test, color='blue', label = 'Actual Value')  
plt.plot(test_data_prediction, color='green', label='Predicted Value')  
plt.title('Actual Price vs Predicted Price')  
plt.xlabel('Number of values')  
plt.ylabel('GLD Price')  
plt.legend()  
plt.show()
```



Score of train and test data

```
regressor.score(X_train,Y_train)
```

0.9984517280531134

```
regressor.score(X_test,Y_test)
```

0.9893081607893282

Making a Predictive System

```
#2671.919922  124.589996  14.0600 15.5100 1.186789
input_data = (2671.919922,  14.0600,  15.5100,  1.186789)

# changing the input_data to a numpy array
input_array = np.asarray(input_data)

# reshape the np array as we are predicting for one instance
input_reshaped = input_array.reshape(1,-1)

test_data_prediction = regressor.predict(input_reshaped)
print(test_data_prediction)

#expected output : 124.589996
```

```
[123.86529702]
```

✓ 0s completed at 4:54 PM

